

Functions and Aggregations

CS 341 Database Systems Lab

One User – One Schema

*A schema is a collection of database objects.
A schema is owned by a database user*

Sqlplus / as sysdba

We are now connected as the admin user of the database which has privileges to create more users.

Setting your user

- SQL> CREATE USER (any user name with prefix c##) IDENTIFIED BY (any password);
- e.g. create user c##myuser identified by 123; *you may create a user with your name or specific to the schema*
- SQL> GRANT UNLIMITED TABLESPACE TO C##MYUSER;
- SQL> GRANT CONNECT, RESOURCE, DBA TO C##MYUSER;

Running Queries via Command Line

Similarly, you can use the SQLPlus Command line

```
Command Prompt - sqlplus

C:\Users\abeeratarig>sqlplus

SQL*Plus: Release 19.0.0.0.0 - Production on Wed Aug 28 22:37:40 2024
Version 19.3.0.0.0

Copyright (c) 1982, 2019, Oracle. All rights reserved.

Enter user-name: c##dblab24
Enter password:
Last Successful login time: Wed Aug 28 2024 22:31:13 +05:00

Connected to:
Oracle Database 19c Enterprise Edition Release 19.0.0.0.0 - Production
Version 19.3.0.0.0

SQL> select * from employees where last_name = 'Jones';

EMPLOYEE_ID FIRST_NAME          LAST_NAME
-----
195 Vance                      Jones
VJONES      650.501.4876      17-MAR-07 SH_CLERK      2800
123                      50

SQL>
```

Recap



Syntax details

- SQL commands are **case insensitive** - SELECT = Select, Product = product
- Values are not, 'Seattle' not equal to 'seattle'
- Use single quotes for constants : 'abc' - best practice (versus "abc" with mixed support)
- To say "don't know the value"/ missing values we use **NULL** E.g., Student GPA in 1st quarter = NULL, not zero
- Free-form language - no special indentation is required but consistent formatting style is recommended for easy maintenance of SQL queries.


SELECT

SELECT c1, c2 FROM t;

Query data in columns c1, c2 from a table

SELECT * FROM t;

Query all rows and columns from a table



Projection is the operation of producing an output table with tuples that have a subset of their prior attributes

WHERE

Selection is the operation of filtering a relation's tuples on some condition

SELECT c1, c2 **FROM** t
WHERE condition;
 Query data and filter rows with a condition

- Comparison Operators →

Operator	Description
=	Equal
>	Greater than
<	Less than
>=	Greater than or equal
<=	Less than or equal
<>	Not equal. Note: In some versions of SQL this operator may be written as !=
BETWEEN	Between a certain range
LIKE	Search for a pattern
IN	To specify multiple possible values for a column

AND, OR, NOT

- Used in **WHERE** clause
- The **AND** operator : all the conditions are TRUE.
- The **OR** operator: if any of the conditions is TRUE.
- The **NOT** operator displays a record if the condition(s) is NOT TRUE.
- Take care of brackets when defining multiple conditions

IN operator

```
SELECT c1, c2 FROM t  
WHERE c1 [NOT] IN value_list;
```

Query rows in a list

- The **IN** operator allows you to specify multiple values in a **WHERE** clause.
- The **IN** operator is a shorthand for multiple **OR** conditions.

Distinct

```
SELECT DISTINCT c1 FROM t  
WHERE condition;
```

Query distinct rows from a table

Between

```
SELECT c1, c2 FROM t  
WHERE c1 BETWEEN low AND high;
```

Query rows between two values

- Selects values within a given range. The values can be numbers, text, or dates.
- The BETWEEN operator is inclusive: begin and end values are included.
- Similar to querying using a combination of \geq and \leq

LIKE (Wildcards)

```
SELECT c1, c2 FROM t1
WHERE c1 [NOT] LIKE pattern;
Query rows using pattern matching %, _
```

- The percent sign (%) represents zero, one, or multiple characters
- The underscore sign (_) represents one, single character

LIKE Operator	Description
WHERE CustomerName LIKE 'a%'	Finds any values that start with "a"
WHERE CustomerName LIKE '%a'	Finds any values that end with "a"
WHERE CustomerName LIKE '%or%'	Finds any values that have "or" in any position
WHERE CustomerName LIKE '_r%'	Finds any values that have "r" in the second position
WHERE CustomerName LIKE 'a_%'	Finds any values that start with "a" and are at least 2 characters in length
WHERE CustomerName LIKE 'a__%'	Finds any values that start with "a" and are at least 3 characters in length
WHERE ContactName LIKE 'a%o'	Finds any values that start with "a" and ends with "o"

Order by

```
SELECT c1, c2 FROM t  
ORDER BY c1 ASC [DESC];
```

Sort the result set in ascending or descending order

- Default ascending order
- DESC - descending
- ASC - ascending

Aliases

Column Name

```
SELECT column_name AS alias_name  
FROM table_name;
```

Table Name

- ```
SELECT column_name(s)
FROM table_name AS alias_name;
```



# Example of Column Aliases

- No spaces in the alias
- If you need spaces then enclose in double quotes

```
SELECT employee_id AS EmployeeID,
 first_name AS FirstName,
 last_name AS LastName
FROM employees;
```

| Query Result x                     |            |           |           |
|------------------------------------|------------|-----------|-----------|
| SQL   Fetched 50 rows in 0 seconds |            |           |           |
|                                    | EMPLOYEEID | FIRSTNAME | LASTNAME  |
| 1                                  | 198        | Donald    | OConnell  |
| 2                                  | 199        | Douglas   | Grant     |
| 3                                  | 200        | Jennifer  | Whalen    |
| 4                                  | 201        | Michael   | Hartstein |
| 5                                  | 202        | Pat       | Fay       |
| 6                                  | 203        | Susan     | Mavris    |
| 7                                  | 204        | Hermann   | Baer      |
| 8                                  | 205        | Shelley   | Higgins   |
| 9                                  | 206        | William   | Gietz     |
| 10                                 | 100        | Steven    | King      |

```
-- ALIAS
SELECT employee_id AS "Employee ID",
 first_name AS "First Name",
 last_name AS "Last Name"
FROM employees;
```

| Query Result x                         |             |            |           |
|----------------------------------------|-------------|------------|-----------|
| SQL   Fetched 50 rows in 0.003 seconds |             |            |           |
|                                        | Employee ID | First Name | Last Name |
| 1                                      | 198         | Donald     | OConnell  |
| 2                                      | 199         | Douglas    | Grant     |
| 3                                      | 200         | Jennifer   | Whalen    |
| 4                                      | 201         | Michael    | Hartstein |

# Aggregation

Lab02A

# Aggregate Functions

- **AVG()** – returns the average of a set.
- **COUNT()** – returns the number of items in a set.
- **MAX()** – returns the maximum value in a set.
- **MIN()** – returns the minimum value in a set
- **SUM()** – returns the sum of all or distinct values in a set

SQL aggregate functions ignore null values.

If using COUNT(\*), it counts all rows including those with NULLs. However, COUNT(column\_name) ignores NULLs in that column.

# Example

```
-- Find the average salary for all employees
select AVG(salary) from employees;
```

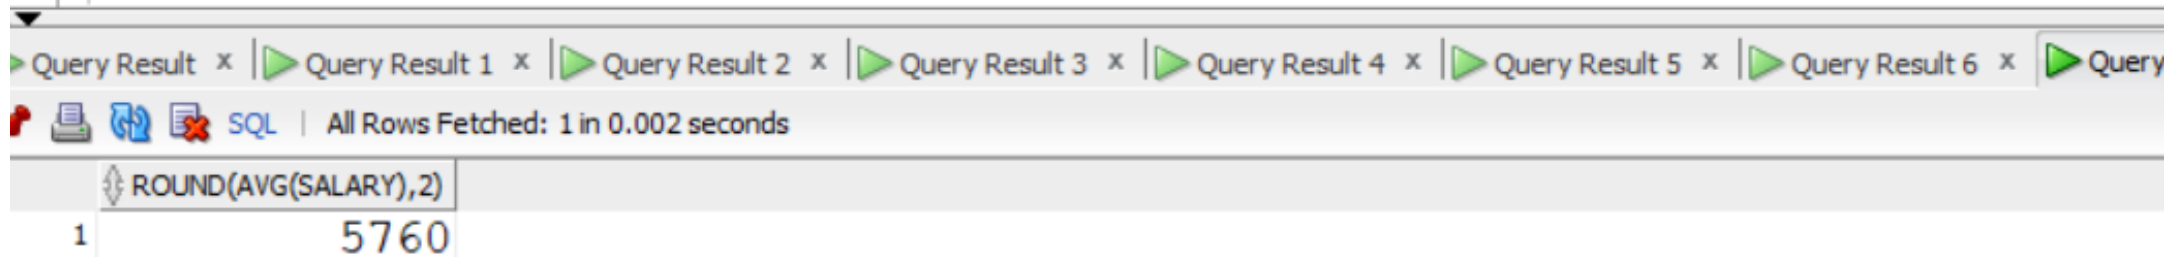
|                                            |                                           |                |                |                |
|--------------------------------------------|-------------------------------------------|----------------|----------------|----------------|
| Query Result                               | Query Result 1                            | Query Result 2 | Query Result 3 | Query Result 4 |
| SQL   All Rows Fetched: 1 in 0.009 seconds |                                           |                |                |                |
| AVG(SALARY)                                |                                           |                |                |                |
| 1                                          | 6461.831775700934579439252336448598130841 |                |                |                |

```
--Round off the salary
select ROUND(AVG(salary),2) from employees;
```

|                                            |                |                |                |                |
|--------------------------------------------|----------------|----------------|----------------|----------------|
| Query Result                               | Query Result 1 | Query Result 2 | Query Result 3 | Query Result 4 |
| SQL   All Rows Fetched: 1 in 0.002 seconds |                |                |                |                |
| ROUND(AVG(SALARY),2)                       |                |                |                |                |
| 1                                          | 6461.83        |                |                |                |

# Avg of specific Job\_id

```
--Find average salary for employees belonging to JOB_ID = 'IT_PROG'
select ROUND(AVG(salary),2) from employees where Job_ID='IT_PROG';
```



The screenshot shows a database query result window with a toolbar at the top containing icons for saving, refreshing, and other actions. Below the toolbar, it indicates 'All Rows Fetched: 1 in 0.002 seconds'. The main area displays a table with one column, 'ROUND(AVG(SALARY),2)', and one row with the value '5760'.

|   | ROUND(AVG(SALARY),2) |
|---|----------------------|
| 1 | 5760                 |

# Avg of all Job\_ids

```
--Find average salary for employees belonging to each JOB_ID
select distinct job_id from employees;
select ROUND(AVG(salary),2) from employees where Job_ID='AD_VP';
select ROUND(AVG(salary),2) from employees where Job_ID='FI_ACCOUNT';
select ROUND(AVG(salary),2) from employees where Job_ID='PU_CLERK';
-- Keep repeating until the end -- Exhaustive, no?
```

| Query Result x | Query Result 1 x                      | Query Result 2 x | Query Result 3 x | Query Result 4 x | Query Result 5 x | Query Result 6 x | Q |
|----------------|---------------------------------------|------------------|------------------|------------------|------------------|------------------|---|
| SQL            | All Rows Fetched: 19 in 0.002 seconds |                  |                  |                  |                  |                  |   |
| JOB_ID         |                                       |                  |                  |                  |                  |                  |   |
| 1 AD VP        |                                       |                  |                  |                  |                  |                  |   |
| 2 FI ACCOUNT   |                                       |                  |                  |                  |                  |                  |   |
| 3 PU CLERK     |                                       |                  |                  |                  |                  |                  |   |
| 4 SH CLERK     |                                       |                  |                  |                  |                  |                  |   |
| 5 HR REP       |                                       |                  |                  |                  |                  |                  |   |
| 6 PU MAN       |                                       |                  |                  |                  |                  |                  |   |
| 7 AC MGR       |                                       |                  |                  |                  |                  |                  |   |
| 8 ST CLERK     |                                       |                  |                  |                  |                  |                  |   |

# Group by

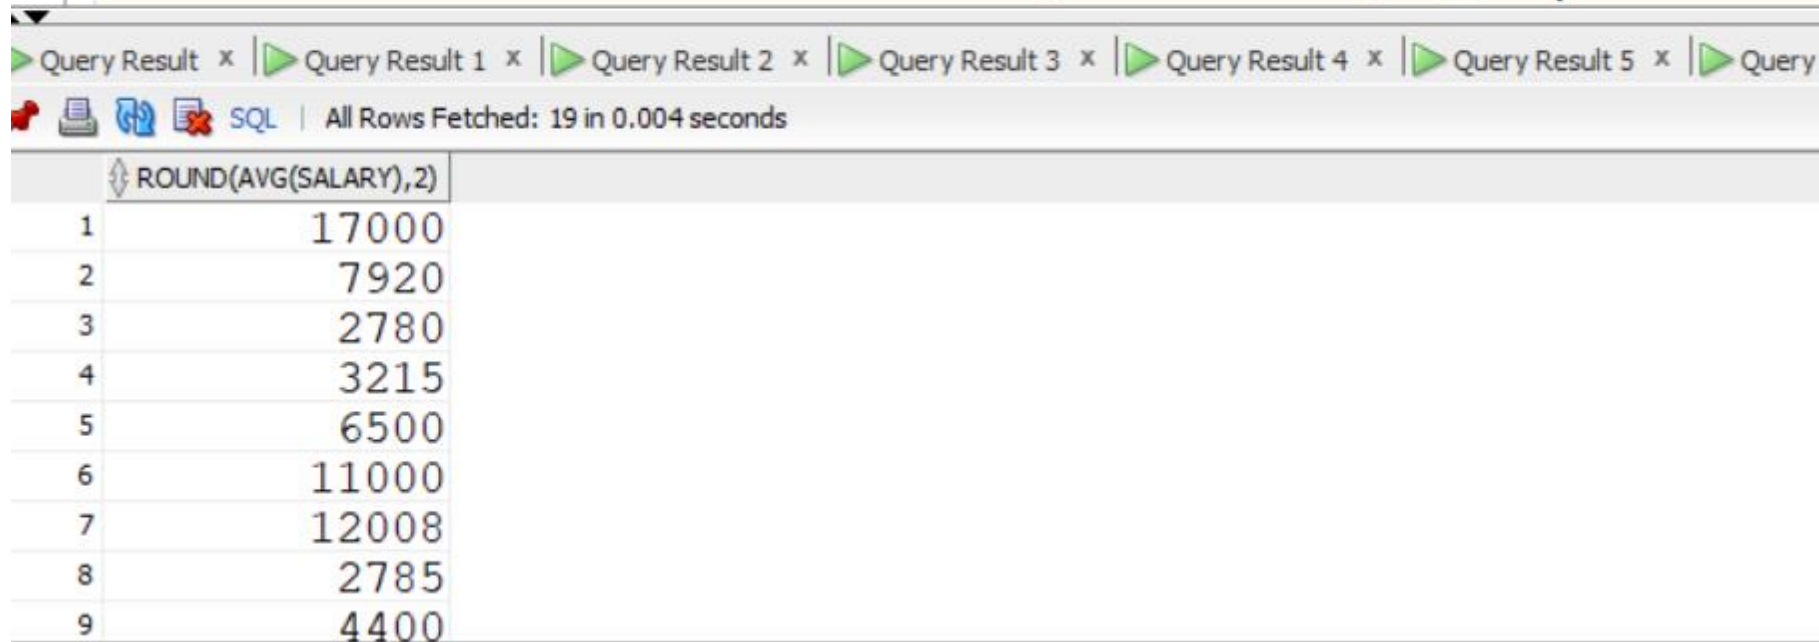
```
SELECT c1, aggregate(c2)
FROM t
GROUP BY c1;
```

Group rows using an aggregate function

- Allows you to group rows based on values of one or more columns. It returns one row for each group.

# Group by Job\_id

```
--use group by job_id
select ROUND(AVG(salary),2) from employees group by job_id;
-- How do I know which value belongs to which job_id?
```



The screenshot shows a database query result window with a tabbed interface. The active tab is 'Query Result', and it displays the results of the SQL query: `select ROUND(AVG(salary),2) from employees group by job_id;`. The window indicates that all 19 rows were fetched in 0.004 seconds. The results are presented in a table with two columns: an implicit index for the job\_id and the calculated average salary rounded to two decimal places.

|   | ROUND(AVG(SALARY),2) |
|---|----------------------|
| 1 | 17000                |
| 2 | 7920                 |
| 3 | 2780                 |
| 4 | 3215                 |
| 5 | 6500                 |
| 6 | 11000                |
| 7 | 12008                |
| 8 | 2785                 |
| 9 | 4400                 |



# Project Job\_id in the results

```
select job_id, ROUND(AVG(salary),2) from employees group by job_id;
```

Query Result x | Query Result 1 x | Query Result 2 x | Query Result 3 x | Query Result 4 x | Query Result 5 x | Query Result 6 x

SQL | All Rows Fetched: 19 in 0.004 seconds

|   | JOB_ID     | ROUND(AVG(SALARY),2) |
|---|------------|----------------------|
| 1 | AD VP      | 17000                |
| 2 | FI ACCOUNT | 7920                 |
| 3 | PU CLERK   | 2780                 |
| 4 | SH CLERK   | 3215                 |
| 5 | HR REP     | 6500                 |
| 6 | PU MAN     | 11000                |
| 7 | AC MGR     | 12008                |
| 8 | ST CLERK   | 2785                 |
| 9 | AD ASST    | 4400                 |

# Adding more columns in SELECT

```
select department_id, job_id, ROUND(AVG(salary),2) from employees group by job_id;
```

Query Result x | Query Result 1 x | Query Result 2 x | Query Result 3 x | Query Result 4 x | Query Result 5 x | Query Result 6 x | Query Result 7 x

SQL | Executing:select department\_id, job\_id, ROUND(AVG(salary),2) from employees group by job\_id in 0 seconds

ORA-00979: not a GROUP BY expression  
00979. 00000 - "not a GROUP BY expression"  
\*Cause:  
\*Action:  
Error at Line: 45 Column: 8

# Group by - Example

| StudentID | Name    | Section | Marks |
|-----------|---------|---------|-------|
| 1         | Alpha   | A       | 60    |
| 2         | Bravo   | B       | 88    |
| 3         | Charlie | A       | 70    |
| 4         | Danny   | B       | 25    |
| 5         | Eric    | B       | 50    |

- Find how many students are there in **each section**

| Section | Count(*) |
|---------|----------|
| A       | 2        |
| B       | 3        |

- `SELECT * FROM Students GROUP BY Section` ❌
- `SELECT Count(*) FROM Students GROUP BY Section`
- `SELECT Section, Count(*) FROM Students GROUP BY Section`
- `SELECT Name, Section, Count(*) FROM Students GROUP BY Section` ❌

# Having

```
SELECT c1, aggregate(c2)
FROM t
GROUP BY c1
HAVING condition;
```

Filter groups using HAVING clause

- Specify a condition for groups, you use the HAVING clause.

The WHERE clause applies the condition to individual rows **before the rows are summarized into groups** by the GROUP BY clause. However, the HAVING clause applies the condition to the groups **after the rows are grouped into groups**.

# Having - Example

| StudentID | Name    | Section | Marks |
|-----------|---------|---------|-------|
| 1         | Alpha   | A       | 60    |
| 2         | Bravo   | B       | 88    |
| 3         | Charlie | A       | 70    |
| 4         | Danny   | B       | 25    |
| 5         | Eric    | B       | 50    |

- `SELECT` Section, Count(\*) *As* NumberOfStudents `FROM` Students `GROUP BY` Section

| Section | NumberOfStudents |
|---------|------------------|
| A       | 2                |
| B       | 3                |

- `SELECT` Section, Count(\*) *As* NumberOfStudents `FROM` Students `GROUP BY` Section `HAVING` Count(\*) < 3

OR

- `SELECT` Section, Count(\*) *As* NumberOfStudents `FROM` Students `GROUP BY` Section `HAVING` NumberOfStudents < 3

| Section | NumberOfStudents |
|---------|------------------|
| A       | 2                |

# Functions

Lab02B

# Numeric Functions

| Functions        | Description                                                                                            |
|------------------|--------------------------------------------------------------------------------------------------------|
| <b>FLOOR</b>     | Returns the largest whole number equal to or less than a specified number.                             |
| <b>MOD</b>       | Returns the modulus of a number.                                                                       |
| <b>POWER</b>     | Returns m_value raised to the n_value power                                                            |
| <b>REMAINDER</b> | Returns the remainder after one numeric expression is divided by another.                              |
| <b>ROUND</b>     | Rounds the number based on the standard rounding rules (e.g., 2.5 rounds to 3, while 2.4 rounds to 2). |
| <b>SQRT</b>      | Computes the square root of an expression.                                                             |
| <b>TRUNC</b>     | Truncates a number to a specified number of decimal places.                                            |

# String Functions





# Upper, Lower, Initcap

- UPPER(expression) → All to uppercase
- LOWER(expression) → All to lowercase
- INITCAP(expression) → Capitalize first character of each word and rest to lower case.

# CONCAT, ||

- Concatenate 2 strings and return the combined string.
- Equivalent to using || operator. (Oracle)
- Other DBMS accept + operator for string concatenation

# TRIM, LTRIM, RTRIM

Remove the space character or other specified characters from the

- TRIM (trim\_source, [set]) → **start or end**
- LTRIM (trim\_source, [set]) → **left** end
- RTRIM (trim\_source, [set]) → **right** end

of a string.

By default, the set will be a single space.

# LPAD, RPAD

- LPAD (source\_string, target\_length [,pad\_string]) → Return a string that is **left-padded** with the specified characters to a certain length.
- RPAD (source\_string, target\_length [,pad\_string]) → Return a string that is **right-padded** with the specified characters to a certain length.
- Default pad\_string is a single space

# REPLACE

Replaces all occurrences of a specified substring in a string with another.

- REPLACE (expression, string\_pattern, string\_replacement)

# SUBSTR

Extract a substring from a string.

- SUBSTR (str, start\_position [, substring\_length, [, occurrence ]])

# INSTR

Search for a substring and return the location of the substring in a string

- INSTR(string , substring [, start\_position [, occurrence]])

# LENGTH

Return the number of characters (or length) of a specified string

- `LENGTH(string_expression);`



# DUAL TABLE

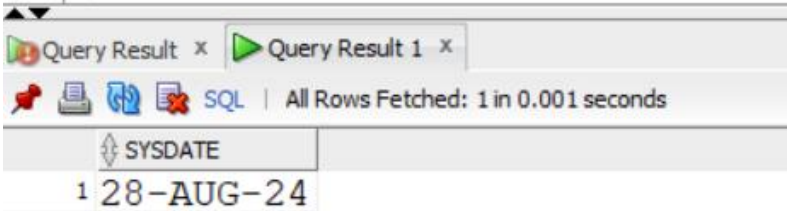
- The DUAL table is owned by the user SYS and can be accessed by users.
- Contains one column, DUMMY, and one row with the value X.
- Useful when you want to return a value once only – for instance, the value of a constant, pseudo-column, or expression that is not derived from a table with user data.

# Date Functions

- **SYSDATE** - Return the current system date and time of the operating system where the Oracle Database resides
- **ADD\_MONTHS** - Add a number of months (n) to a date and return the same day which is n of months away.
- **MONTHS\_BETWEEN** - Return the number of months between two dates.

# Example

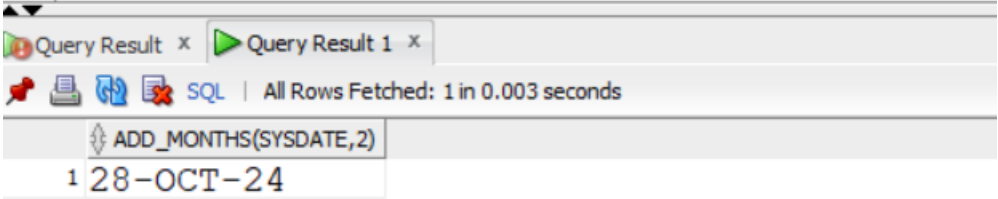
```
select SYSDATE from dual;
```



The screenshot shows a SQL query result window with the following details:

- Query: `select SYSDATE from dual;`
- Execution time: 1 in 0.001 seconds
- Column: `SYSDATE`
- Result: `1 28-AUG-24`

```
-- ADDMONTHS
select ADD_MONTHS(SYSDATE,2) from dual;
```



The screenshot shows a SQL query result window with the following details:

- Query: `-- ADDMONTHS`  
`select ADD_MONTHS(SYSDATE,2) from dual;`
- Execution time: 1 in 0.003 seconds
- Column: `ADD_MONTHS(SYSDATE,2)`
- Result: `1 28-OCT-24`

Wednesday, August 28

August 2024

| Su | Mo | Tu | We | Th | Fr | Sa |
|----|----|----|----|----|----|----|
| 28 | 29 | 30 | 31 | 1  | 2  | 3  |
| 4  | 5  | 6  | 7  | 8  | 9  | 10 |
| 11 | 12 | 13 | 14 | 15 | 16 | 17 |
| 18 | 19 | 20 | 21 | 22 | 23 | 24 |
| 25 | 26 | 27 | 28 | 29 | 30 | 31 |
| 1  | 2  | 3  | 4  | 5  | 6  | 7  |

# Example

```
--MONTHS_BETWEEN Date1 - Date 2
```

```
select MONTHS_BETWEEN(SYSDATE, '1-OCT-23') from dual;
```

Query Result 7 x

SQL | All Rows Fetched: 1 in 0.001 seconds

|   | MONTHS_BETWEEN(SYSDATE,'1-OCT-23')        |
|---|-------------------------------------------|
| 1 | 10.90080458482676224611708482676224611708 |

```
-- ROUND MONTHS
```

```
select ROUND(MONTHS_BETWEEN(SYSDATE, '1-OCT-23')) from dual;
```

Query Result x

SQL | All Rows Fetched: 1 in 0.002 seconds

|   | ROUND(MONTHS_BETWEEN(SYSDATE,'1-OCT-23')) |
|---|-------------------------------------------|
| 1 | 11                                        |

# Date Functions

- **NEXT\_DAY** - Get the first weekday that is later than a specified date.
- **LAST\_DAY** - Gets the last day of the month of a specified date.

# Example

```
--NEXT_DAY, Get the next saturday after today
select NEXT_DAY(SYSDATE, 'Saturday') from dual;
```

| Query Result x                      |  |
|-------------------------------------|--|
| All Rows Fetched: 1 in 0.01 seconds |  |
| NEXT_DAY(SYSDATE, 'SATURDAY')       |  |
| 1 31-AUG-24                         |  |

```
--LAST_DAY, get the last day of this month
select LAST_DAY(SYSDATE) from dual;
```

| Query Result x                       |  |
|--------------------------------------|--|
| All Rows Fetched: 1 in 0.001 seconds |  |
| LAST_DAY(SYSDATE)                    |  |
| 1 31-AUG-24                          |  |

# Date Functions

- **ROUND** - Return a date rounded to a specific unit of measure.
- **TRUNC** - Return a date truncated to a specific unit of measure.

# Example

```
--ROUND - Return a date rounded to a specific unit of measure.- Execute to see output
select ROUND (DATE '2024-08-24', 'DD') from dual;
select ROUND (DATE '2024-08-24', 'MM') from dual;
select ROUND (DATE '2024-08-24', 'YY') from dual;

-- TRUNC - Return a date truncated to a specific unit of measure.
select TRUNC (DATE '2024-08-24', 'MM') from dual;
```

| Query Result x                             |           |
|--------------------------------------------|-----------|
| SQL   All Rows Fetched: 1 in 0.001 seconds |           |
| TRUNC (DATE '2024-08-24', 'MM')            |           |
| 1                                          | 01-AUG-24 |